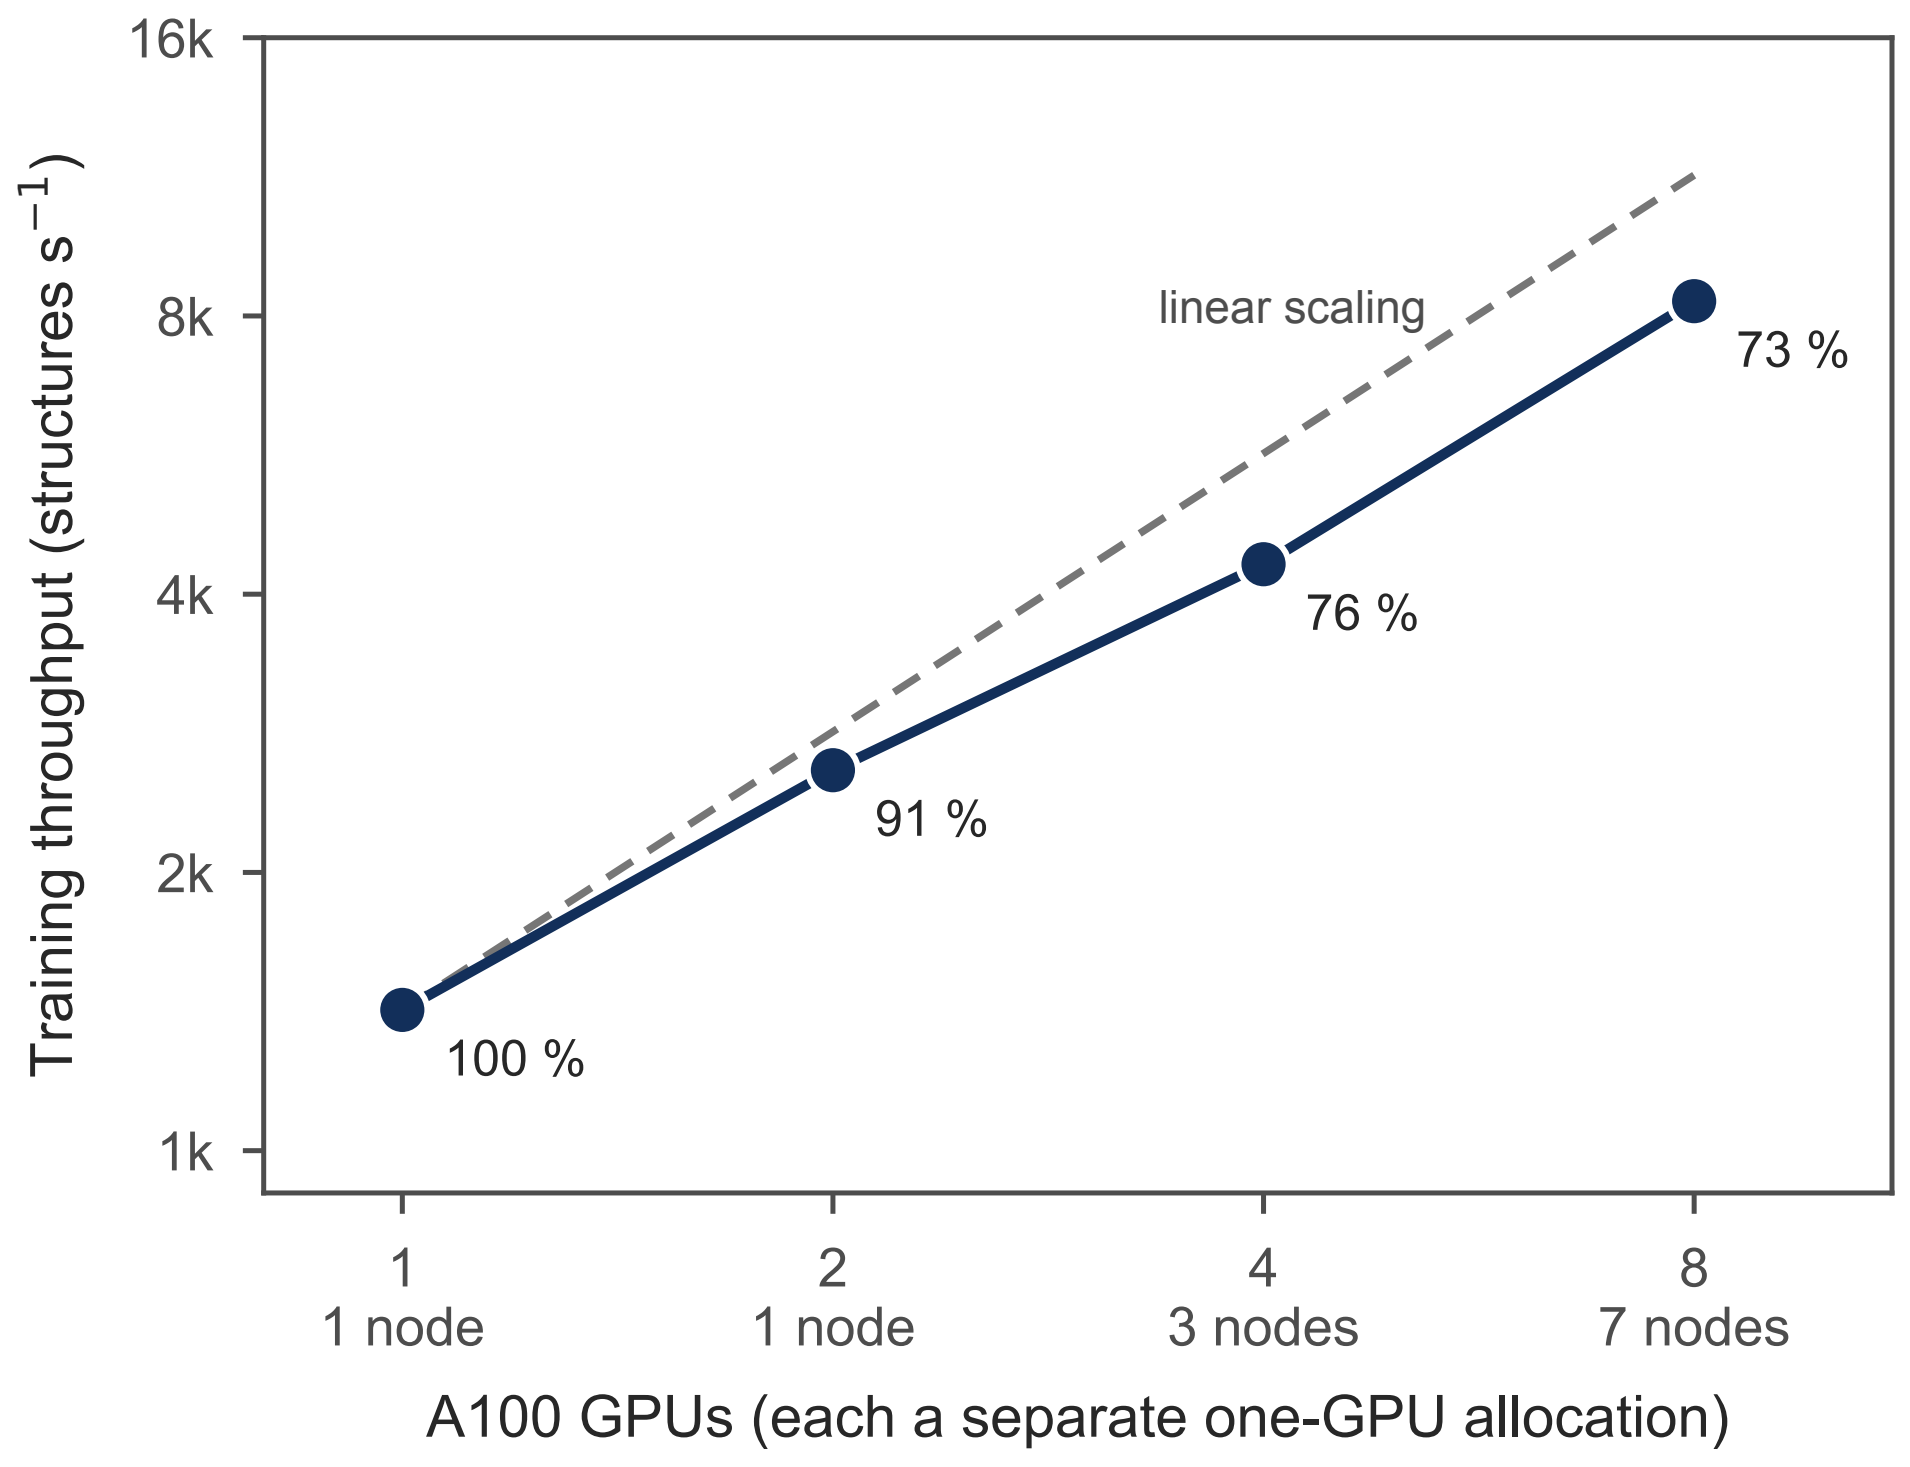
16k
8k
4k
2k
1k
Training throughput (structures s^{-1})
linear scaling
73 %
76 %
91 %
100 %
1
1 node
2
1 node
4
3 nodes
8
7 nodes
A100 GPUs (each a separate one-GPU allocation)